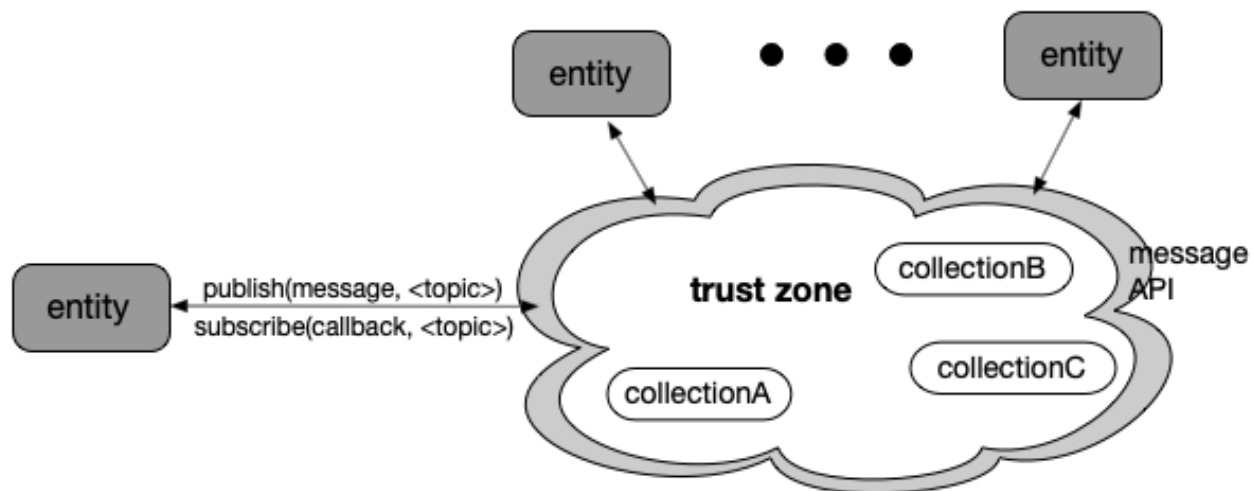


Home IoT Example using Message-Based Publish/Subscribe (mbps)

The Data-Centric Toolkit's modules can be used to create bespoke transports that are integrated with a Versec trust schema. The examples in this directory use a message-based pub/sub shim that presents applications with an MQTT-inspired interface. This directory contains three example applications, with a home IoT flavor and three versions of its trust schema. The DCT "shim", **mbps.hpp**, implements an mbps client and presents an API based on *messages* as an application level unit of information. Unlike its MQTT inspiration, mbps is *brokerless*, its clients connecting instead to a (fully distributed) *trust zone*. Three different levels of communication security are supplied in the three example trust schemas: `iot{0-2}.trust`. The applications remain the same (same binary). When started, individual application entities are given an *identity bundle* containing the trust anchor, the trust schema, and their own signing identity. It is this identity bundle that determines the particular trust zone the entity joins. Use of an mbps client lets entities publish and subscribe to messages in a particular trust zone; all mechanics of the trust zone are transparent to an application which only "sees" a message space.



Messages for a target may be further specified with subtopics and subscriptions can be particularized by these or may just subscribe to all messages (in the mbps shim's sync protocol collection). An mbps client is a client of a *trust zone* not of a *server* (or broker), and its implementation reflects that. For example, in MQTT, a message of type CONNECT is the first packet sent from the client to the server and MQTT applications typically call a client *connect* method, passing in a callback function that handles a CONNACK message from the server. Since mbps is serverless, its *connect* method is used to do **all** initialization required in order to communicate as a member of the trust zone (acquire keys, etc), then calls back to an application handler once communications become possible.

Build notes

These examples work on Macs and Linux and uses IPv6 multicast. They use the Data-Centric Toolkit (<https://github.com/pollere/DCT>) which currently uses the ndn-ind library (<https://github.com/operantnetworks>) and, for broadcast performance, NFD version with the the *no-nacks-on-multicast-faces* and *ship-pending-interests-on-register* NFD patches as well as ndn-cxx patch (for keys) available at: <https://github.com/pollere/NDNpatches>.

[Also must fix the bug that causes rejection of registrations within the same millisecond before can use key distributors.] After cloning the DCT repo, the example programs can be used. Do a 'make' in your examples/mbps directory.

Using the examples

This **hmlot** directory contains three applications (app1.cpp, app2.cpp, and app3.cpp) and three different text file trust schemas: `iot{0,1,2}.trust`. The example applications are meant to provide simple illustrations of the use of mbps, DCT, and Versec trust schemas. Each trust schema implements a different security model that can be used with any of the applications interchangeably. The `iot1` and `iot2` trust schemas use a signing chain where publications are signed by a role cert that is signed by the root cert. The `iot0` trust schema checks that `syncData` and `Publications` are signed by a valid cert (same trust root) and follow a basic format, but doesn't enforce any particular role cert to issue certain publications (in contrast, `iot1.trust` and `iot2.trust` only let *operators* issue commands). In `iot1`, publications are signed with their signing keys using the EdDSA signature manager (DCT's cert distributor automatically publishes its own signing key chain during connect and acquires and validate that of others continuously) and Data packets are signed by their sender using the EdDSA signature manager. The `iot2` trust schema requires the AEAD signature manager for Data packets so those will be encrypted (DCT's group key distributor is invoked automatically during connect to make and securely distribute those keys) while the `Publications` use the EdDSA signature manager. In a DCT-based transport, entities **must** publish their signing chains before they can communicate ("join the trust zone").

In the example applications, *operator* roles subscribe to all messages in the application domain and device roles only subscribe to the **command** topic and locations matching their IDs and the location designator *all*. A single handler function (`msgRecv`) is used for all subscriptions. The `connect` lambda sets up subscriptions, creates the first message body and invokes `msgPubr` which uses the client's publish function and schedules another message to be sent. The mbps client's `run()` is called after `connect()`. The files also have examples of using message confirmation (roughly similar to MQTT QoS 1).

Creating Identity Bundles

Identity bundles of certificates are required to run these examples. The bundles contain, in order, the trust root certificate, the trust schema to be used (in cert form and signed by the trust root), and the signing chain for the entity. Bundles do **not** contain the signing certs of other devices that will be part of the application domain; the trust schema, common root-of-trust, and run-time library handle this, allowing entities to discover new members dynamically. The trust schema can be compiled using `schemaCompile` (recommend installation in `tools` directory). The `DCT/tools` directory supplies utilities to create certs and id bundles and an application is invoked directly with its identity bundle.

Acting as a Configurer, first, use the `schemaCompile` utility in the `versec` directory to compile the text file of trust rules into a binary schema, for `iotn`, e.g.:

```
schemaCompile -o iot1.scm iot1.trust
```

Next, a trust anchor (root) cert for this network and domain is created as a self-signed cert. The `-s` option is for the publication signature manager, the `-o` names the output file and finally the publication prefix for this application domain is listed:

```
make_cert -s EdDSA -o iot1.root iot1
```

The output file contains a cert with the public key as content and the associated signing (private) key. The next step is to make a cert carrying the schema binary signed by the root cert in accordance with rules in the trust schema:

```
schema_cert -o iot1.schema iot1.scm iot1.root
```

The cert content is the trust schema binary so there is no signing key in the output file. Next, make identity certs that are signed by the trust anchor. Make one for each planned instance of the application, particularizing by role (here **operator** or **device**) and by specifying an id within that role, e.g.:

```
make_cert -s EdDSA -o alice.cert iot1/operator/alice iot1.root
```

Make an identity bundle for each of these certs using the `make_bundle` utility in the `tools` directory, this one is given the identity name, e.g.:

```
make_bundle -o alice.bundle iot1.root iot1.schema +alice.cert
```

The "+" on the role cert indicates that its signing (private) key should be included in the bundle. The other certs in the bundle don't (and shouldn't) have signing keys. In order to run the applications, at least two identities are needed, so make at least two bundles. In general, more identity bundles can be made later for additional application instances. A shell script `mkIDs.sh` is included that will generate five example bundles for the trust rule file given on the command line.

Running the examples

The host must be running an NDN Forwarding Daemon (with Pollere's patches) and must have a localnet route to the IPv6 multicast Face. At least two application entities should be started or nothing will happen. (For now, do a "nfdc cs erase /localnet" before starting a new session.) `app1` publishes the given number of messages, sending a short message each time. If the `-p` option is set, the application will persist waiting for more messages.

Example usage:

```
nfdc cs erase /localnet
```

```
app1 -c 10 <identity>.bundle or app1 -c 10 -p <identity>.bundle
```

`app2` publishes an initial message and waits to receive a message before scheduling another message publication at a short delay.

Example usage:

```
app2 <identity>.bundle
```

`app3` randomizes some publication components and removes the confirmation callbacks of `app2`.

Example usage:

```
app3 <identity>.bundle
```

In addition to a message body, mbps.hpp's *msgParms* struct is used to pass additional information about a message with the MBPS client that is needed by the MBPS transport. These parameters are required by the trust schema to fill in some of the *tags* defined for Publications. Looking at the supplied trust schemas, you will see which tags need to be set by the application (no leading "_" and not set in mbps) and that the applications supply them, e.g.:

```
msgParms{{"target", capability},{ "topic", "event"s},{ "trgtLoc",myId},{ "topicArgs",  
myState}}
```

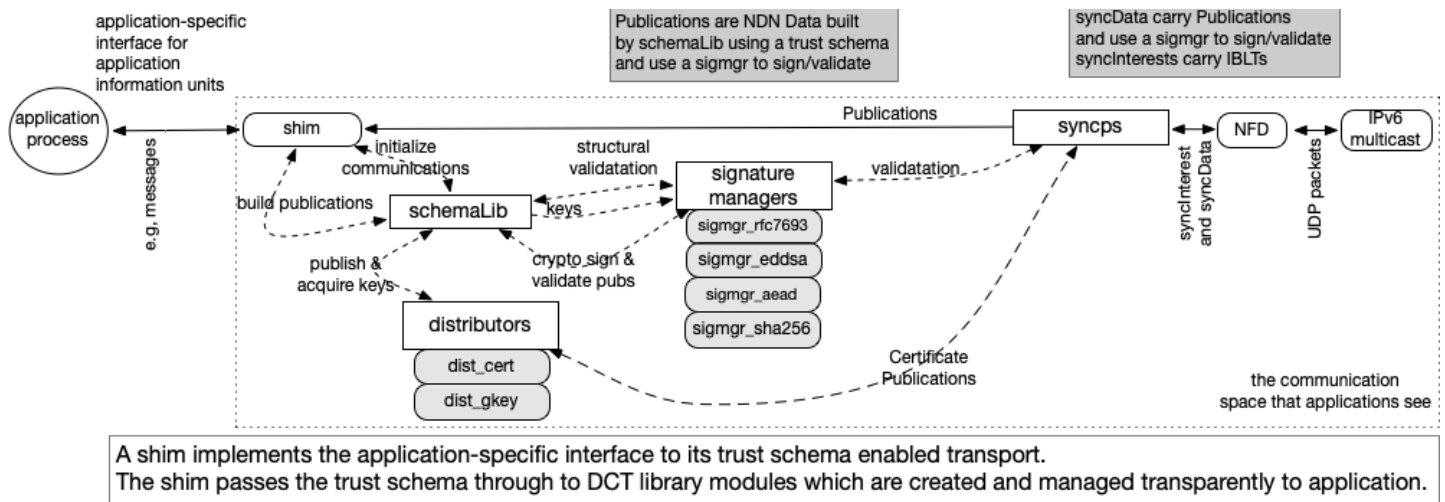
Tags can be accessed in a received message from the mbpsMsg structure and the applications show this in their msgRecv() methods.

Elements of a DCT Bespoke Transport

This discussion contains more information about how the mbps shim uses the DCT library than is necessary in order to simply write applications and trust schemas.

mbps's shim serves as an API and passes an identity bundle to modules from Pollere's Data-Centric Toolkit (DCT) which implement the corresponding trust zone. (DCT currently uses a version of the NDN Forwarding Daemon NFD to provide network layer primitives over an IPv6 multicast network.) This section provides an overview of the module interaction. This is an area of active work, that is, we are still refining how the "pieces" go together; this is the current snapshot. The modules that uniquely define a particular transport are the shim (here DCT/examples/shims/mbps.hpp) and the trust schema (e.g., *iot0.trust*) while the rest are DCT library modules. The transport translates application messages into secured Publications (and Publications to messages) that are used by its *syncps* Sync protocol that manages the set synchronization of a particular Collection of Publications using NDN primitives (*syncInterests* and *syncData*). Publications are NDN Data which have distinct name components (trust schema tags) and a data portion that holds the message or a segment of a large message. (Publications are encapsulated in NDN Sync Data, which are exchanged with NDN forwarders, through *syncps*. Since the focus is on broadcast communications, those go to the **localnet** Face.)

DCT's library is currently separated into four directories that cover different functionalities: *syncps* (the *syncPubSub* object and its helpers), *sigmgrs* (sign and validate Data using different algorithms), *distributors* (maintain distributed cert or key collections), and *schema* (holds the *schemaLib* modules). The figure shows how the various modules of DCT are used by a custom shim (like *mbps.hpp*) to make a bespoke transport that presents an application-specific interface and uses NFD and its IPv6 multicast interface to send and receive packets on a local network. In addition to the cryptographic signing and validation of Data packets (e.g., RFC7693, EdDSA, SHA256, AEAD), the *schemaLib* is employed to do construction and structural validation of Publications, ensuring that all of the schema's naming constraints hold for the publication and its signing chain. DCT does not permit trust schemas to specify integrity signing (RFC7693 or SHA256) for publications as this is at odds with DCT's security model, but these are used within *distributors* to sign *syncData* during initialization.



MBPS Publication Names

Each syncps handles the synchronization of Publications within its collection, which may encompass an entire application domain or may be further specified by a longer prefix. The transport used in these examples uses a collection for the publications that carry application messages (wirePrefix/trustSchemaTP/pubs), a collection for the signing cert chains (wirePrefix/trustSchemaTP/cert), and, if AEAD is specified for wire validation, a collection for the symmetric group keys (wirePrefix/trustSchemaTP/keys). The latter two are transparent to the application. For application messages, publication names are constructed according to the following table, where the component structure can be used for hierarchical topic and subtopic handling and the individual components are specified in, and can be secured through, the domain's trust schema. The last three components are required by mbps and will be filled in by the shim. The rest depend on the particular trust schema, both in the number of components (tags from the trust schema definition) and whether they must be supplied by the application through the mbps API. Any component (tag) in the publication definition(s) in the trust schema that does not have a leading "_" and is not one of the final three components must be supplied.

component	description
networkID	This is a single component in our examples, denoting the domain of iot {1,2,3}. Multiple components might be needed for some deployments.
target	a capability, attribute, property
topic	gives the type of message (command, report)
trgtLoc	target location id or alias, e.g. device(s), room, local, component (smarththings)
topicArgs	more specific information for this topic
origin	identifies the publisher: role-specific ID and deviceID
msgID	unsigned 32 bit number to uniquely identify each message
sCnt	unsigned 16 bit integer; indicates if single publication message or multiples (see below)
mts	time mbps client receives message from application process in UTC microsecs

A target is the "audience" for the Publication. This can be used in many ways, but one (IoT-centric) way to think of it is as a *capability* e.g., lock, alarm, colorControl, motionSensor. In general, topic can be used in many ways; in these examples it indicates the type of message carried, e.g. an event, a command, or a status. The target location can be a particular device, class of device, component of a device, a room, a vicinity. The arguments can supply any additional information associated with that topic. The origin component can be used to identify the publisher (only used for debugging), msgID is set to a hash in mbps.hpp. sCnt is packaged in an unsigned 16 bit integer which, if not 0 to indicate a single publication message, uses the upper 8 bits for the sequence of this piece within the total message and the lower 8 bits for the total number of pieces in the message. This means that any message must fit into 255 publication segments (so byte-limited to maximum content size times 255) but the current shims limit the number of segments to 64 for an efficient implementation and for the iblt implementation, the limit should be kept under 80.

Confirmation of Publications ("QoS")

MBPS provides mechanisms that enable applications to set a callback (if desired) indicating whether its message made it into the collection. This relies on mechanisms in syncps that can check that its Publication was seen on a SyncInterest issued by another process and methods in the shim that ensure all the Publications of any particular message have been confirmed. This gives an indication of message publication "success" roughly analogous to MQTT's QoS 1. It is invoked at the application by passing a callback to the mbps.hpp publish() method as the third argument. If the callback argument is not used, the option is not employed by the client. (Although mbps may or may not use it for publications, the mechanism is employed for initial signing key distribution in the connect method, transparently to the application.)

Notes on the included trust schemas

Identity bundles can be made using each of these trust schemas. Since a trust schema determines a particular trust zone, applications invoked with identity bundles with different trust schema will not communicate with each other. The security model can be changed just by using a set of identity bundles with the new trust schema while the application binary remains the same.

iot0.trust

An entity needs the trust schema and a valid signing identity in order to sign and validate Publications. Both Publications and syncData are signed with the private key of the signing cert in the identity bundle using the EdDSA signature manager. There is no role differentiation so the schema has only type of Publication.

Management and distribution of the signing certs is handled within the schema library modules and is transparent to both the application and mbps.hpp. The distributor *dist_cert.hpp* manages access to a cert Collection that contains all the certs needed to validate Publication and sync Data for an application. During connect(), each entity publishes its signing cert chain and subscribes to that of others in the Collection; the application is called back when signing chain publication is confirmed.

iot1.trust

This schema introduces operator and device role certs and derives additional Publication types that can require particular signing certs to be considered valid. As in iot0.trust, the EdDSA signature manager is used for cryptographic signing of both Publications and syncData.

iot2.trust

This schema differs from iot1.trust only in that the wireValidator ensures privacy by encrypting all sync Data payloads using dynamically created and distributed symmetric keys. Its AEAD signature manager requires a group key so an additional distributor, *dist_gkey.hpp* is automatically employed. Any client whose role permits it to create and distribute group keys attempts to become the key maker for the domain. (The simple self-selection algorithm used here could certainly be improved upon but is provided as an illustration and starting point.) Once a valid symmetric key is obtained, messages can be encrypted/decrypted so the application connect handler is called back. The AEAD signature manager is used to encrypt/decrypt SyncData while the Publications use EdDSA.